

Large Language Models

CSC413 Tutorial 9

Yongchao Zhou

Overview

- What are LLMs?
- Why LLMs?
- Emergent Capabilities
 - Few-shot In-context Learning
 - Advanced Prompt Techniques
- LLM Training
 - Architectures
 - Objectives
- LLM Finetuning
 - Instruction finetuning
 - RLHF
 - Bootstrapping
- LLM Risks

What are Language Models?

- Narrow Sense
 - A probabilistic model that assigns a probability to every finite sequence (grammatical or not)

Sentence: “the cat sat on the mat”

$$P(\text{the cat sat on the mat}) = P(\text{the}) * P(\text{cat}|\text{the}) * P(\text{sat}|\text{the cat})$$

$$*P(\text{on}|\text{the cat sat}) * P(\text{the}|\text{the cat sat on})$$


$$*P(\text{mat}|\text{the cat sat on the})$$

Implicit order




- Broad Sense
 - Decoder-only models (GPT-X, OPT, LLaMA, PaLM)
 - Encoder-only models (BERT, RoBERTa, ELECTRA)
 - Encoder-decoder models (T5, BART)

Large Language Models - **Billions of Parameters**




Large Language Models - **Hundreds of Billions of Tokens**



Large Language Models - **yottaFlops of Compute**

AI training runs, estimated computing resources used


Floating-point operations, selected systems, by type, log scale



Why LLMs?

- **Scaling Law for Neural Language Models**


- Performance depends strongly on scale! We keep getting better performance as we scale the model, data, and compute up!



Why LLMs?

- **Generalization**

- We can now use one single model to solve many NLP tasks




<https://arxiv.org/pdf/1910.10683.pdf>

Why LLMs?

- **Emergent Abilities**

- Some ability of LM is not present in smaller models but is present in larger models



https://docs.google.com/presentation/d/1yzbmYB5E7G8IY2-KzhmArmPYwwl7o7CUST1xRZDUu1Y/edit?resourcekey=0-6_TnUMoKWCK_FN2BiPxmbw#slide=id.g1fc34b3ac18_0_27

Emergent Capability - In-Context Learning

Traditional fine-tuning (not used for GPT-3)


Fine-tuning

The model is trained via repeated gradient updates using a large corpus of example tasks.




One-shot

In addition to the task description, the model sees a single example of the task. No gradient updates are performed.



Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.




<https://arxiv.org/pdf/2005.14165.pdf>


Emergent Capability - In-Context Learning

	No Prompt	Prompt
Zero-shot (os)	skicts = sticks	Please unscramble the letters into a word, and write that word: skicts = sticks
1-shot (1s)	chiar = chair skicts = sticks	Please unscramble the letters into a word, and write that word: chiar = chair skicts = sticks
Few-shot (FS)	chiar = chair [...] pciinc = picnic skicts = sticks	Please unscramble the letters into a word, and write that word: chiar = chair [...] pciinc = picnic skicts = sticks


Emergent Capability - In-Context Learning



Pretraining + Fine-tuning Paradigm



Pre-training



Fine-Tuning

Pretraining + Prompting Paradigm

- Fine-tuning (FT)
 - + Strongest performance
 - - Need curated and labeled dataset for each new task (typically 1k-100k ex.)
 - - Poor generalization, spurious feature exploitation
- Few-shot (FS)
 - + Much less task-specific data needed
 - + No spurious feature exploitation
 - - Challenging
- One-shot (1S)
 - + "Most natural," e.g. giving humans instructions
 - - Challenging
- Zero-shot (OS)
 - + Most convenient
 - - Challenging, can be ambiguous

**Stronger
task-specific
performance**



**More convenient,
general, less data**

Emergent Capability - Chain of Thoughts Prompting

Standard Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27. X

Chain-of-Thought Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?



A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. ✓

Emergent Capability - Chain of Thoughts Prompting



Emergent Capability - Zero Shot CoT Prompting

(a) Few-shot

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A:

(Output) The answer is 8. X

(b) Few-shot-CoT

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A:

(Output) *The juggler can juggle 16 balls. Half of the balls are golf balls. So there are $16 / 2 = 8$ golf balls. Half of the golf balls are blue. So there are $8 / 2 = 4$ blue golf balls. The answer is 4. ✓*

(c) Zero-shot

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A: The answer (arabic numerals) is

(Output) 8 X


(d) Zero-shot-CoT (Ours)

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?


A: **Let's think step by step.**

(Output) *There are 16 balls in total. Half of the balls are golf balls. That means that there are 8 golf balls. Half of the golf balls are blue. That means that there are 4 blue golf balls. ✓*


Emergent Capability - Zero Shot CoT Prompting



Emergent Capability - Self-Consistency Prompting



Emergent Capability - Least-to-Most Prompting



Emergent Capability - Augmented Prompting Abilities

Advanced Prompting Techniques

- Zero-shot CoT Prompting
- Self-Consistency
- Divide-and-Conquer


Ask a human to

- Explain the rationale
- Double check the answer
- Decompose to easy subproblems


Large Language Models demonstrate some human-like behaviors!

Training Architectures

Encoder-decoder models (T5, BART)



Decoder-only models (GPT-X, PaLM)




Training Objectives - UL2

R-Denoising	S-Denoising	X-Denoising
<p>Inputs:</p> <p>[R] He dealt in archetypes before anyone knew such things existed, and his 3 to take an emotion or a situation 5 it to the limit helped create a cadre of plays that have been endlessly 4 – and copied. Apart from this, Romeo and Juliet inspired Malorie Blackman's Noughts 5 there are references to Hamlet in Lunar Park by Bret Easton Ellis 2 The Tempest was the cue for The Magus by John Fowles.</p> <p>Target:</p>	<p>Inputs:</p> <p>[S] He dealt in archetypes before anyone knew such things existed, and his ability to take an emotion or a situation and push it to the limit helped create a cadre of plays that have been endlessly staged – and copied. Apart from this, Romeo and Juliet 95</p> <p>Target:</p>	<p>Inputs:</p> <p>[X] He dealt in archetypes 16 things existed, and his ability to take an emotion or a situation 32 plays that have been endlessly staged – and copied. Apart from 24 Malorie Blackman's Noughts & Crosses, there are references to Hamlet in Lunar 24 Tempest was the cue for The Magus by John Fowles.</p> <p>Target:</p>
		<p>Inputs:</p> <p>[X] He dealt in archetypes 3 anyone knew such things existed, ai 3 ability to take an 5 situation and push it to the limit helped 4 cadre of plays 4 been endlessly staged – and 5 Apart from this, Romeo and Juliet inspired Malorie Blackman's 5 Crosses, 3 are references to Hamlet in 3 Park by Bret Easton 2 and 4 4 was the 2 for The 4 by John 5</p> <p>Target:</p>

What kinds of things does pretraining learn?

- *Stanford University is located in _____, California.* [Trivia]
- *I put __ fork down on the table.* [syntax]
- *The woman walked across the street, checking for traffic over __ shoulder.* [coreference]
- *I went to the ocean to see the fish, turtles, seals, and ____.* [lexical semantics/topic]
- *Overall, the value I got from the two hours watching it was the sum total of the popcorn and the drink. The movie was ____.* [sentiment]
- Iroh went into the kitchen to make some tea. Standing next to Iroh, Zuko pondered his destiny. Zuko left the _____. [some reasoning – this is harder]
- I was thinking about the sequence that goes 1, 1, 2, 3, 5, 8, 13, 21, ____ [some basic arithmetic; they don't learn the Fibonnaci sequence]

Finetune - Instruction Finetune



Finetune - Instruction Finetune

Finetuning tasks

TO-SF

Commonsense reasoning
Question generation
Closed-book QA
Adversarial QA
Extractive QA
Title/context generation
Topic classification
Struct-to-text
...

*55 Datasets, 14 Categories,
193 Tasks*

Muffin

Natural language inference
Code instruction gen.
Program synthesis
Dialog context generation
Closed-book QA
Conversational QA
Code repair
...

69 Datasets, 27 Categories, 80 Tasks

CoT (Reasoning)

Arithmetic reasoning Explanation generation
Commonsense Reasoning Sentence composition
Implicit reasoning ...

9 Datasets, 1 Category, 9 Tasks


Natural Instructions v2

Cause effect classification
Commonsense reasoning
Named entity recognition
Toxic language detection
Question answering
Question generation
Program execution
Text categorization
...



*372 Datasets, 108 Categories,
1554 Tasks*

- ❖ A Dataset is an original data source (e.g. SQuAD).
- ❖ A Task Category is unique task setup (e.g. the SQuAD dataset is configurable for multiple task categories such as extractive question answering, query generation, and context generation).
- ❖ A Task is a unique <dataset, task category> pair, with any number of templates which preserve the task category (e.g. query generation on the SQuAD dataset.)

Finetune - Instruction Finetune



Finetune - Instruction Finetune



Finetune - RLHF

Step 1

Collect demonstration data, and train a supervised policy.


A prompt is sampled from our prompt dataset.



A labeler demonstrates the desired output behavior.



This data is used to fine-tune GPT-3 with supervised learning.



Step 2

Collect comparison data, and train a reward model.

A prompt and several model outputs are sampled.



A labeler ranks the outputs from best to worst.



This data is used to train our reward model.



Step 3

Optimize a policy against the reward model using reinforcement learning.


A new prompt is sampled from the dataset.



The policy generates an output.



The reward model calculates a reward for the output.



The reward is used to update the policy using PPO.

r_k

Application - ChatGPT

ChatGPT



Examples

"Explain quantum computing in simple terms" →



Capabilities

"Got any creative ideas for a 10 year old's birthday?" →

Allows user to provide follow-up corrections



Limitations

May occasionally generate incorrect information


May occasionally produce harmful instructions or biased content

"How do I make an HTTP request in Javascript?" →

Trained to decline inappropriate requests


Limited knowledge of world and events after 2021

Application - ChatGPT



<https://yaofu.notion.site/How-does-GPT-Obtain-its-Ability-Tracing-Emergent-Abilities-of-Language-Models-to-their-Source-s-b9a57ac0fcf74f30a1ab9e3e36fa1dc1>

Finetune - Bootstrapping




Q: What can be used to carry a small dog?

Answer Choices:

- (a) swimming pool
- (b) basket
- (c) dog show
- (d) backyard
- (e) own home

A: The answer must be something that can be used to carry a small dog. Baskets are designed to hold things. Therefore, the answer is basket (b).

Finetune - Bootstrapping



Large Language models Risks

- LLMs make mistakes
(falsehoods, hallucinations)
- LLMs can be misused
(misinformation, spam)
- LLMs can cause harms
(toxicity, biases, stereotypes)
- LLMs can be attacked
(adversarial examples, poisoning, prompt injection)
- LLMs can be useful as defenses
(content moderation, explanations)



The image shows a news article from CNET. The header features the CNET logo in red and the tagline "Your guide to a better future". Below the header, the category "Tech > Services & Software" is listed. The main title of the article is "It's Scary Easy to Use ChatGPT to Write Phishing Emails", displayed in large, bold, black text.

Large language models associate Muslims with violence


[Abubakar Abid](#), [Maheen Farooqi](#) & [James Zou](#) 

[Nature Machine Intelligence](#) 3, 461–463 (2021) | [Cite this article](#)



Resources for further reading

- <https://web.stanford.edu/class/cs224n/>
- <https://stanford-cs324.github.io/winter2022/>
- <https://stanford-cs324.github.io/winter2023/>
- <https://www.cs.princeton.edu/courses/archive/fall22/cos597G/>
- <https://rycolab.io/classes/llm-s23/>
- <https://yaofu.notion.site/How-does-GPT-Obtain-its-Ability-Tracing-Emergent-Abilities-of-Language-Models-to-their-Sources-b9a57ac0fcf74f30a1ab9e3e36fa1dc1>
- <https://www.jasonwei.net/blog/emergence>


Emergent Capability - In-Context Learning



Emergent Capability - Decomposed Prompting




Training Objectives - UL2




Training Techniques - **Parallelism**


Data Parallelism




Pipeline Parallelism



Tensor Parallelism



Expert Parallelism



An illustration of various parallelism strategies on a three-layer model. Each color refers to one layer and dashed lines separate different GPUs.

Training Techniques - **Parallelism**

